

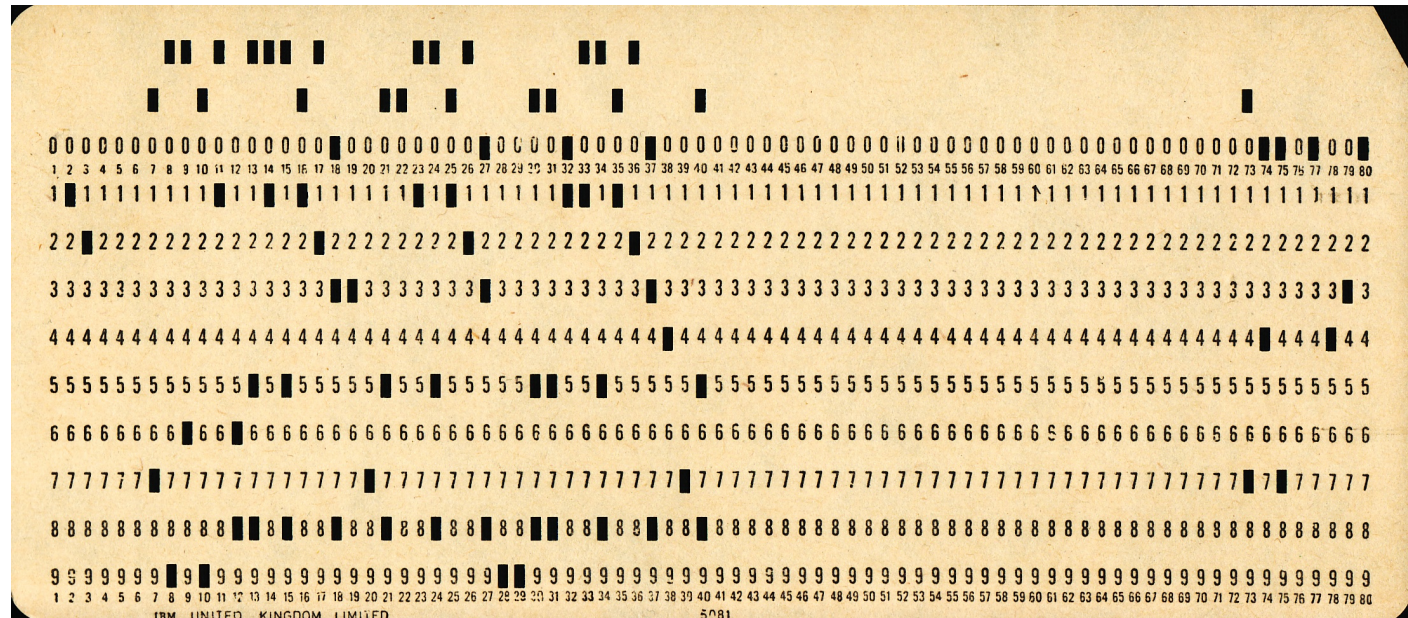
# Day 1, Part 3: Pandas and Data I/O

Introduction to Python

Tom Paskhalis

RECSM Summer School 2023

# Rectangular data



History of rectangular data goes back to punchcards with origins in US census data processing.

Source: [Wikipedia](#)

# Tidy data

- Tidy data is a specific subset of rectangular data, where:
  - Each variable is in a column
  - Each observation is in a row
  - Each value is in a cell

country	year	cases	population
Afghanistan	1999	1815	19987071
Afghanistan	2000	2566	20095360
Brazil	1999	37737	172006362
Brazil	2000	84488	174004898
China	1999	212258	1272915272
China	2000	213766	1286425583

variables

country	year	cases	population
Afghanistan	1999	1815	19987071
Afghanistan	2000	2566	20095360
Brazil	1999	37737	172006362
Brazil	2000	84488	174004898
China	1999	212258	1272915272
China	2000	213766	1286425583

observations

country	year	cases	population
Afghanistan	1999	1815	19987071
Afghanistan	2000	2566	20095360
Brazil	1999	37737	172006362
Brazil	2000	84488	174004898
China	1999	212258	1272915272
China	2000	213766	1286425583

values

Source: [R for Data Science](#)

# Data in Python

- Python can hold and manipulate  $> 1$  dataset at the same time
- Python stores objects in memory
- The limit on the size of data is determined by your computer memory
- Most functionality for dealing with data is provided by external libraries

# Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

# Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

```
In [1]: # Representing 3x3 matrix with list  
mat = [[1, 2, 3],  
        [4, 5, 6],  
        [7, 8, 9]]
```

# Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

```
In [1]: # Representing 3x3 matrix with list  
mat = [[1, 2, 3],  
        [4, 5, 6],  
        [7, 8, 9]]
```

```
In [2]: # Subsetting 2nd row, 3rd element  
mat[1][2]
```

```
Out[2]: 6
```

# Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

```
In [1]: # Representing 3x3 matrix with list  
mat = [[1, 2, 3],  
        [4, 5, 6],  
        [7, 8, 9]]
```

```
In [2]: # Subsetting 2nd row, 3rd element  
mat[1][2]
```

```
Out[2]: 6
```

```
In [3]: # Naturally, this representation  
# breaks down rather quickly  
mat * 2
```

```
Out[3]: [[1, 2, 3], [4, 5, 6], [7, 8, 9], [1, 2, 3], [4, 5, 6], [7, 8,  
9]]
```



# NumPy - numerical analysis in Python

- NumPy (**N**umeric **P**ython) package provides the basis of numerical computing in Python:
  - multidimensional array
  - mathematical functions for arrays
  - array data I/O
  - linear algebra, RNG, FFT, ...

# NumPy - numerical analysis in Python

- NumPy (**N**umeric **P**ython) package provides the basis of numerical computing in Python:
  - multidimensional array
  - mathematical functions for arrays
  - array data I/O
  - linear algebra, RNG, FFT, ...

```
In [4]: # Using 'as' allows to avoid typing full name  
# each time the module is referred to  
import numpy as np
```

# NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

# NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

```
In [5]: arr = np.array([[1, 2, 3],  
                        [4, 5, 6],  
                        [7, 8, 9]])
```

# NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

```
In [5]: arr = np.array([[1, 2, 3],  
                        [4, 5, 6],  
                        [7, 8, 9]])
```

```
In [6]: arr[1][2]
```

```
Out[6]: 6
```

# NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

```
In [5]: arr = np.array([[1, 2, 3],  
                        [4, 5, 6],  
                        [7, 8, 9]])
```

```
In [6]: arr[1][2]
```

```
Out[6]: 6
```

```
In [7]: arr * 2
```

```
Out[7]: array([[ 2,  4,  6],  
              [ 8, 10, 12],  
              [14, 16, 18]])
```

# Working with arrays

# Working with arrays

```
In [8]: # Object type  
        type(arr)
```

```
Out[8]: numpy.ndarray
```



# Working with arrays

```
In [8]: # Object type  
type(arr)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: # Array dimensionality  
arr.ndim
```

```
Out[9]: 2
```

# Working with arrays

```
In [8]: # Object type  
type(arr)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: # Array dimensionality  
arr.ndim
```

```
Out[9]: 2
```

```
In [10]: # Array size  
arr.shape
```

```
Out[10]: (3, 3)
```

# Working with arrays

```
In [8]: # Object type  
type(arr)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: # Array dimensionality  
arr.ndim
```

```
Out[9]: 2
```

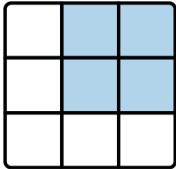
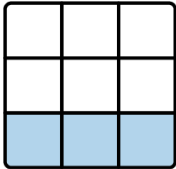
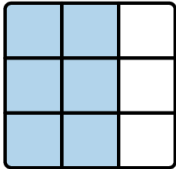
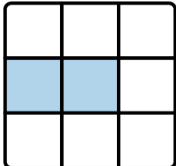
```
In [10]: # Array size  
arr.shape
```

```
Out[10]: (3, 3)
```

```
In [11]: # Calculating summary statistics on array  
# axis indicates the dimension  
# compare to R's `apply(arr, 1, mean)`  
# note that every list within a list  
# is treated as a column (not row)  
arr.mean(axis = 0)
```

```
Out[11]: array([4., 5., 6.])
```

# Array indexing and slicing

	Expression	Shape
	<code>arr[:2,1:]</code>	<code>(2,2)</code>
	<code>arr[2]</code> <code>arr[2, :]</code> <code>arr[2:, :]</code>	<code>(3,)</code> <code>(3,)</code> <code>(1,3)</code>
	<code>arr[:, :2]</code>	<code>(3,2)</code>
	<code>arr[1, :2]</code> <code>arr[1:2, :2]</code>	<code>(2,)</code> <code>(1,2)</code>

Source: [Python for Data Analysis](#)

# Pandas

- Standard Python library does not have data type for tabular data
- However, `pandas` library has become the de facto standard for data manipulation
- `pandas` is built upon (and often used in conjunction with) other computational libraries
- E.g. `numpy` (array data type), `scipy` (linear algebra) and `scikit-learn` (machine learning)

# Pandas

- Standard Python library does not have data type for tabular data
- However, `pandas` library has become the de facto standard for data manipulation
- `pandas` is built upon (and often used in conjunction with) other computational libraries
- E.g. `numpy` (array data type), `scipy` (linear algebra) and `scikit-learn` (machine learning)

In [12]: *# Using 'as' allows to avoid typing full name each time the module is imported*  
`import pandas as pd`

# Series

- *Series* is a one-dimensional array-like object

# Series

- *Series* is a one-dimensional array-like object

```
In [13]: sr1 = pd.Series([150.0, 120.0, 3000.0])  
sr1
```

```
Out[13]: 0      150.0  
         1      120.0  
         2     3000.0  
         dtype: float64
```



# Series

- *Series* is a one-dimensional array-like object

```
In [13]: sr1 = pd.Series([150.0, 120.0, 3000.0])  
sr1
```

```
Out[13]: 0      150.0  
         1      120.0  
         2     3000.0  
         dtype: float64
```

```
In [14]: sr1[0] # Slicing is similar to standard Python objects
```

```
Out[14]: 150.0
```

# Series

- *Series* is a one-dimensional array-like object

```
In [13]: sr1 = pd.Series([150.0, 120.0, 3000.0])  
sr1
```

```
Out[13]: 0    150.0  
         1    120.0  
         2   3000.0  
         dtype: float64
```

```
In [14]: sr1[0] # Slicing is similar to standard Python objects
```

```
Out[14]: 150.0
```

```
In [15]: sr1[sr1 > 200]
```

```
Out[15]: 2    3000.0  
         dtype: float64
```

# Indexing in Series

- Another way to think about Series is as a ordered dictionary

# Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

# Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

```
In [17]: sr2 = pd.Series(d)  
sr2
```

```
Out[17]: apple      150.0  
         banana     120.0  
         watermelon 3000.0  
         dtype: float64
```

# Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

```
In [17]: sr2 = pd.Series(d)
sr2
```

```
Out[17]: apple      150.0
         banana     120.0
         watermelon 3000.0
         dtype: float64
```

```
In [18]: sr2[0] # Recall that this slicing would be impossible for standard dict
```

```
Out[18]: 150.0
```

# Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

```
In [17]: sr2 = pd.Series(d)
sr2
```

```
Out[17]: apple      150.0
         banana     120.0
         watermelon 3000.0
         dtype: float64
```

```
In [18]: sr2[0] # Recall that this slicing would be impossible for standard dict
```

```
Out[18]: 150.0
```

```
In [19]: sr2.index
```

```
Out[19]: Index(['apple', 'banana', 'watermelon'], dtype='object')
```

# DataFrame - the workhorse of data analysis

- *DataFrame* is a rectangular table of data



# DataFrame - the workhorse of data analysis

- *DataFrame* is a rectangular table of data

```
In [20]: data = {'fruit': ['apple', 'banana', 'watermelon'], # DataFrame can be
                'weight': [150.0, 120.0, 3000.0],           # a dict of equal-l
                'berry': [False, True, True]}
df = pd.DataFrame(data)
df
```

```
Out[20]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True
2	watermelon	3000.0	True

# Indexing in DataFrame

- DataFrame has both row and column indices
- `DataFrame.loc()` provides method for *label* location
- `DataFrame.iloc()` provides method for *index* location

# Indexing in DataFrame

- DataFrame has both row and column indices
- `DataFrame.loc()` provides method for *label* location
- `DataFrame.iloc()` provides method for *index* location

```
In [21]: df.iloc[0] # First row
```

```
Out[21]: fruit      apple  
weight    150.0  
berry      False  
Name: 0, dtype: object
```

# Indexing in DataFrame

- DataFrame has both row and column indices
- `DataFrame.loc()` provides method for *label* location
- `DataFrame.iloc()` provides method for *index* location

```
In [21]: df.iloc[0] # First row
```

```
Out[21]: fruit      apple  
weight    150.0  
berry      False  
Name: 0, dtype: object
```

```
In [22]: df.iloc[:,0] # First column
```

```
Out[22]: 0      apple  
1      banana  
2  watermelon  
Name: fruit, dtype: object
```

# Summary of indexing in DataFrame

Expression	Selection Operation
<code>df[val]</code>	Column or sequence of columns +convenience (e.g. slice)
<code>df.loc[lab_i]</code>	Row or subset of rows by label
<code>df.loc[:, lab_j]</code>	Column or subset of columns by label
<code>df.loc[lab_i, lab_j]</code>	Both rows and columns by label
<code>df.iloc[i]</code>	Row or subset of rows by integer position
<code>df.iloc[:, j]</code>	Column or subset of columns by integer position
<code>df.iloc[i, j]</code>	Both rows and columns by integer position
<code>df.at[lab_i, lab_j]</code>	Single scalar value by row and column label
<code>df.iat[i, j]</code>	Single scalar value by row and column integer position

Extra: [Pandas documentation on indexing](#)

# Subsetting in DataFrame



# Subsetting in DataFrame

In [23]: `df.iloc[:2]` *# Select the first two rows (with convenience shortcut for*

Out[23]:

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True





# Subsetting in DataFrame

```
In [23]: df.iloc[:2] # Select the first two rows (with convenience shortcut for
```

```
Out[23]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [24]: df[:2] # Shortcut
```

```
Out[24]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True



# Subsetting in DataFrame

```
In [23]: df.iloc[:2] # Select the first two rows (with convenience shortcut for
```

```
Out[23]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [24]: df[:2] # Shortcut
```

```
Out[24]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [25]: df.loc[:, ['fruit', 'berry']] # Select the columns 'fruit' and 'berry'
```

```
Out[25]:
```

	fruit	berry
0	apple	False
1	banana	True
2	watermelon	True



# Subsetting in DataFrame

```
In [23]: df.iloc[:2] # Select the first two rows (with convenience shortcut for
```

```
Out[23]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [24]: df[:2] # Shortcut
```

```
Out[24]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [25]: df.loc[:, ['fruit', 'berry']] # Select the columns 'fruit' and 'berry'
```

```
Out[25]:
```

	fruit	berry
0	apple	False
1	banana	True
2	watermelon	True

```
In [26]: df[['fruit', 'berry']] # Shortcut
```

Out[26]:

	<b>fruit</b>	<b>berry</b>
<b>0</b>	apple	False
<b>1</b>	banana	True
<b>2</b>	watermelon	True

# Columns in DataFrame



# Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

# Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

```
In [28]: df.columns[0] # This Index object is subsettable
```

```
Out[28]: 'fruit'
```

# Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

```
In [28]: df.columns[0] # This Index object is subsettable
```

```
Out[28]: 'fruit'
```

```
In [29]: df.columns.str.startswith('fr') # As column names are strings, we can a
```

```
Out[29]: array([ True, False, False])
```

# Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

```
In [28]: df.columns[0] # This Index object is subtable
```

```
Out[28]: 'fruit'
```

```
In [29]: df.columns.str.startswith('fr') # As column names are strings, we can a
```

```
Out[29]: array([ True, False, False])
```

```
In [30]: df.iloc[:,df.columns.str.startswith('fr')] # This is helpful with more
```

```
Out[30]:
```

	fruit
0	apple
1	banana
2	watermelon

# Filtering in DataFrame

# Filtering in DataFrame

```
In [31]: df[df.loc[:, 'berry'] == False] # Select rows where fruits are not berry
```

```
Out[31]:
```

	fruit	weight	berry
0	apple	150.0	False

# Filtering in DataFrame

```
In [31]: df[df.loc[:, 'berry'] == False] # Select rows where fruits are not berries
```

```
Out[31]:
```

	fruit	weight	berry
0	apple	150.0	False

```
In [32]: df[df['berry'] == False] # The same can be achieved with more concise syntax
```

```
Out[32]:
```

	fruit	weight	berry
0	apple	150.0	False

# Filtering in DataFrame

```
In [31]: df[df.loc[:, 'berry'] == False] # Select rows where fruits are not berry
```

```
Out[31]:
```

	fruit	weight	berry
0	apple	150.0	False

```
In [32]: df[df['berry'] == False] # The same can be achieved with more concise s
```

```
Out[32]:
```

	fruit	weight	berry
0	apple	150.0	False

```
In [33]: weight200 = df[df['weight'] > 200] # Create new dataset with rows where  
weight200
```

```
Out[33]:
```

	fruit	weight	berry
2	watermelon	3000.0	True



# Variable transformation

- Lambda functions can be used to transform data with `map( )` method

# Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [34]: df['fruit'].map(lambda x: x.upper())
```

```
Out[34]: 0      APPLE  
         1     BANANA  
         2  WATERMELON  
         Name: fruit, dtype: object
```

# Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [34]: df['fruit'].map(lambda x: x.upper())
```

```
Out[34]: 0      APPLE  
         1     BANANA  
         2  WATERMELON  
         Name: fruit, dtype: object
```

```
In [35]: transform = lambda x: x.capitalize()
```

# Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [34]: df['fruit'].map(lambda x: x.upper())
```

```
Out[34]: 0      APPLE  
         1     BANANA  
         2  WATERMELON  
         Name: fruit, dtype: object
```

```
In [35]: transform = lambda x: x.capitalize()
```

```
In [36]: transformed = df['fruit'].map(transform)
```

# Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [34]: df['fruit'].map(lambda x: x.upper())
```

```
Out[34]: 0      APPLE
         1     BANANA
         2  WATERMELON
         Name: fruit, dtype: object
```

```
In [35]: transform = lambda x: x.capitalize()
```

```
In [36]: transformed = df['fruit'].map(transform)
```

```
In [37]: transformed
```

```
Out[37]: 0      Apple
         1     Banana
         2  Watermelon
         Name: fruit, dtype: object
```

# File object

- File object in Python provides the main interface to external files
- In contrast to other core types, file objects are created not with a literal,
- But with a function, `open ( )` :

```
<variable_name> = open(<filepath>, <mode>)
```

# Data input and output

- Modes of file objects allow to:
  - ( `r` )ead a file (default)
  - ( `w` )rite an object to a file
  - e( `x` )clusively create, failing if a file exists
  - ( `a` )ppend to a file
- You can `r+` mode if you need to read and write to file

## Data output example



## Data output example

```
In [38]: f = open('../temp/test.txt', 'w') # Create a new file object in write m
```

## Data output example

```
In [38]: f = open('../temp/test.txt', 'w') # Create a new file object in write mode
```

```
In [39]: f.write('This is a test file.') # Write a string of characters to it
```

```
Out[39]: 20
```

## Data output example

```
In [38]: f = open('../temp/test.txt', 'w') # Create a new file object in write mode
```

```
In [39]: f.write('This is a test file.') # Write a string of characters to it
```

```
Out[39]: 20
```

```
In [40]: f.close() # Flush output buffers to disk and close the connection
```

# Data input example

- To avoid keeping track of open file connections, `with` statement can be used

Extra: [Python documentation on with statement](#)

# Data input example

- To avoid keeping track of open file connections, `with` statement can be used

Extra: [Python documentation on with statement](#)

```
In [41]: with open('../temp/test.txt', 'r') as f: # Note that we use 'r' mode for reading
          text = f.read()
```

# Data input example

- To avoid keeping track of open file connections, `with` statement can be used

Extra: [Python documentation on with statement](#)

```
In [41]: with open('../temp/test.txt', 'r') as f: # Note that we use 'r' mode for reading
          text = f.read()
```

```
In [42]: text
```

```
Out[42]: 'This is a test file.'
```

# Reading and writing data in pandas

- pandas provides high-level methods that takes care of file connections
- These methods all follow the same read\_<format> and to\_<format> name patterns
- CSV (comma-separated value) files are the standard of interoperability

```
<variable_name> = pd.read_<format>(<filepath>)
```

```
<variable_name>.to_<format>(<filepath>)
```

# Reading data in `pandas` example

- We will use the data from [Kaggle 2021 Machine Learning and Data Science Survey](#)
- For more information you can read the [executive summary](#)
- Or explore the [winning Python Jupyter Notebooks](#)



## Reading data in `pandas` example

- We will use the data from [Kaggle 2021 Machine Learning and Data Science Survey](#)
- For more information you can read the [executive summary](#)
- Or explore the [winning Python Jupyter Notebooks](#)

```
In [43]: # We specify that we want to combine first two rows as a header  
kaggle2021 = pd.read_csv('../data/kaggle_survey_2021_responses.csv', he
```

```
/tmp/ipykernel_279893/1791299071.py:2: DtypeWarning: Columns (19  
5,201) have mixed types. Specify dtype option on import or set l  
ow_memory=False.  
    kaggle2021 = pd.read_csv('../data/kaggle_survey_2021_response  
s.csv', header = [0,1])
```

# Visual data inspection



# Visual data inspection

```
In [44]: kaggle2021.head() # Returns the top n (n=5 default) rows
```

```
Out[44]:
```

	Time from Start to Finish (seconds)	Q1	Q2	Q3	Q4	Q5	
	Duration (in seconds)	What is your age (# years)?	What is your gender? - Selected Choice	In which country do you currently reside?	What is the highest level of formal education that you have attained or plan to attain within the next 2 years?	Select the title most similar to your current role (or most recent title if retired): - Selected Choice	For how years you writing program
0	910	50-54	Man	India	Bachelor's	Other	5-10

					degree		
<b>1</b>	784	50-54	Man	Indonesia	Master's degree	Program/Project Manager	20+
<b>2</b>	924	22-24	Man	Pakistan	Master's degree	Software Engineer	1-3
<b>3</b>	575	45-49	Man	Mexico	Doctoral degree	Research Scientist	20+
<b>4</b>	781	45-49	Man	India	Doctoral degree	Other	< 1

5 rows × 369 columns

Visual data inspection continued



# Visual data inspection continued

In [45]: `kaggle2021.tail()` # Returns the bottom n (n=5 default) rows

Out[45]:

	Time from Start to Finish (seconds)	Q1	Q2	Q3	Q4	Q5		
	Duration (in seconds)	What is your age (# years)?	What is your gender? - Selected Choice	In which country do you currently reside?	What is the highest level of formal education that you have attained or plan to attain within the next 2 years?	Select the title most similar to your current role (or most recent title if retired): - Selected Choice	For how r years you l writing (a ar programm	
	25968	1756	30-34	Man	Egypt	Bachelor's	Data	1-3 y



					degree	Analyst	
<b>25969</b>	253	22-24	Man	China	Master's degree	Student	1-3 y
<b>25970</b>	494	50-54	Man	Sweden	Doctoral degree	Research Scientist	I have r written
<b>25971</b>	277	45-49	Man	United States of America	Master's degree	Data Scientist	5-10 y
<b>25972</b>	255	18-21	Man	India	Bachelor's degree	Business Analyst	I have r written

5 rows × 369 columns

## Reading in other (non-`.csv`) data files

- Pandas can read in file other than `.csv` (comma-separated value)
- Common cases include STATA `.dta`, SPSS `.sav` and SAS `.sas`
- Use `pd.read_stata(path)`, `pd.read_spss(path)` and `pd.read_sas(path)`
- Check [here](#) for more examples

## Writing data out in pandas

- Note that when writing data out we start with the object name storing the dataset
- I.e. `df.to_csv(path)` as opposed to `df = pd.read_csv(path)`
- Pandas can also write out into other data formats
- E.g. `df.to_excel(path)`, `df.to_stata(path)`

## Writing data out in pandas

- Note that when writing data out we start with the object name storing the dataset
- I.e. `df.to_csv(path)` as opposed to `df = pd.read_csv(path)`
- Pandas can also write out into other data formats
- E.g. `df.to_excel(path)`, `df.to_stata(path)`

```
In [46]: kaggle2021.to_csv('../temp/kaggle2021.csv')
```

# Additional pandas materials

Books:

- McKinney, Wes. 2022. *Python for Data Analysis: Data Wrangling with pandas, NumPy, and Jupyter*. 3rd ed. Sebastopol, CA: O'Reilly Media

**From the original author of the library!**

Online:

- [Pandas Getting Started Tutorials](#)
- [Pandas Documentation](#) (intermediate and advanced)

# Tomorrow

- Exploratory data analysis
- Data visualization

